

Evaluating the Feasibility of Sampling-Based Techniques for Training Multilayer Perceptrons

Sana Ebrahimi

Rishi Advani

Abolfazl Asudeh

University of Illinois Chicago
{sebrah7, radvani2, asudeh}@uic.edu

28th International Conference on Extending Database Technology
March 28, 2025 – Barcelona, Spain



Outline

- 1 Motivation
- 2 Overview
- 3 ALSH-APPROX
- 4 MC-APPROX
- 5 Theoretical Analysis
- 6 Highlighted Experiments
- 7 Final Remarks

A scalability issue

- DNNs are often very large in scale:
 - ▶ Computationally expensive to train
 - ▶ Requiring powerful hardware, including expensive GPUs
 - ▶ CPUs are widely available! why not using them instead?

Motivation

A scalability issue

- DNNs are often very large in scale:
 - ▶ Computationally expensive to train
 - ▶ Requiring powerful hardware, including expensive GPUs
 - ▶ CPUs are widely available! why not using them instead?

Extended Database Technologies for Machine Learning

- Feasibility of the extension of sampling-based techniques for efficient training of deep neural networks (DNNs)
- On CPU machines with limited resources

Motivation

A scalability issue

- DNNs are often very large in scale:
 - ▶ Computationally expensive to train
 - ▶ Requiring powerful hardware, including expensive GPUs
 - ▶ CPUs are widely available! why not using them instead?

Extended Database Technologies for Machine Learning

- Feasibility of the extension of sampling-based techniques for efficient training of deep neural networks (DNNs)
- On CPU machines with limited resources

Two scalability directions


- ① Sampling-based approaches based on locality-sensitive hashing (LSH) [GIM99]
- ② Monte-carlo (MC) estimations [Rob16]

Outline

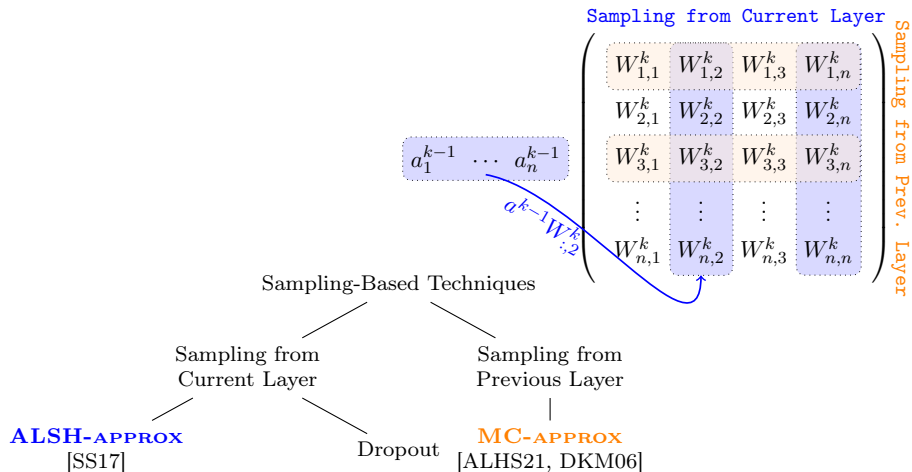
- 1 Motivation
- 2 Overview
- 3 ALSH-APPROX
- 4 MC-APPROX
- 5 Theoretical Analysis
- 6 Highlighted Experiments
- 7 Final Remarks

Review: multi-layer perceptron (MLP)

$$z^k = a^{k-1}W^k + b^k$$
$$a^k = f(z^k)$$



Taxonomy of Sampling-Based Techniques



Outline

- 1 Motivation
- 2 Overview
- 3 ALSH-APPROX**
- 4 MC-APPROX
- 5 Theoretical Analysis
- 6 Highlighted Experiments
- 7 Final Remarks

Asymmetric Locality-Sensitive Hashing

A family \mathcal{H} is $(\tau, c, p_1 > p_2)$ -sensitive for c -NNS, if for all $h \in \mathcal{H}$:


$$\begin{aligned} \text{sim}(w, a) \geq \tau &\implies \Pr[h(Q(a)) = h(P(w))] \geq p_1 \\ \text{sim}(w, a) \leq c\tau &\implies \Pr[h(Q(a)) = h(P(w))] \leq p_2 \end{aligned}$$

For $w, a \in \mathbb{R}^n$ with $\|w\| \leq C$, where C is a constant less than 1, and $\|a\| = 1$, they define the transformations P and Q for the inner product as follows.

$$\begin{aligned} P: \mathbb{R}^n &\rightarrow \mathbb{R}^{n+m}, & w &\mapsto [w; \|w\|^{2^1}, \dots, \|w\|^{2^m}] \\ Q: \mathbb{R}^n &\rightarrow \mathbb{R}^{n+m}, & a &\mapsto [a; 1/2, \dots, 1/2] \end{aligned} \tag{1}$$

ALSH-APPROX

- Only computes the high-value elements of $a^{k-1}W^k$ (called active nodes), skipping the computation for the small values.



Uses ALSH to find active nodes: the columns $W_{:,j}^k$ that collide with a^{k-1} into the same bin.

$$\operatorname{argmax}_j \langle W_{:,j}^k, a^{k-1} \rangle \approx \operatorname{argmin}_j \|Q(a^{k-1}) - P(W_{:,j}^k)\|$$

Outline

- 1 Motivation
- 2 Overview
- 3 ALSH-APPROX
- 4 MC-APPROX**
- 5 Theoretical Analysis
- 6 Highlighted Experiments
- 7 Final Remarks

Approximate Matrix Multiplication $A \times B$

- $AB_{i,j} = \sum_{t=1}^n A_{i,t}B_{t,j}$
- **Goal:** A Monte-carlo method to estimate the sum over a subset of values of t
- **Standard approach:** uniformly select a sample of columns (high estimation error)
- **Idea** (weighted sampling): compute the values that are expected to be larger


$$p_t = \frac{\|A_{:,t}\| \cdot \|B_{t,:}\|}{\sum_{k=1}^n \|A_{:,k}\| \cdot \|B_{k,:}\|}$$

For a set C of c samples:

$$AB_{i,j} \approx \sum_{t \in C} \frac{1}{cp_t} A_{i,t} B_{t,j}$$

MC-APPROX

- Computes all values of $a^{k-1}W^k$, but instead of exact computation, it **estimates** each value.
- It is more **appropriate for mini-batch settings**, where a matrix of activation vectors is multiplied to W^k .
- For single-row settings, it may not accurately estimate p_t values \Rightarrow its estimations are not accurate.



Outline

- 1 Motivation
- 2 Overview
- 3 ALSH-APPROX
- 4 MC-APPROX
- 5 Theoretical Analysis**
- 6 Highlighted Experiments
- 7 Final Remarks

Negative Result: ALSH-APPROX does not scale

Theorem

Let f be a linear activation function such that $a = f(z) = z$. Suppose for any node n_p^l , $\sum_{i \in \uparrow_p^l} a_i^{l-1} W_{i,p} = c \sum_{i \notin \uparrow_p^l} a_i^{l-1} W_{i,p}$. Then, $a_j^k = \bar{a}_j^k \left(\frac{c+1}{c} \right)^k$. That is,


$$e_j^k = \bar{a}_j^k \left(\left(\frac{c+1}{c} \right)^k - 1 \right)$$

Example: suppose $c = 5$ (i.e., the weighted sum for the active nodes is five times that of the inactive nodes). Then,

k (number of layers)	1	2	3	4	5	6
e_j^k / \bar{a}_j^k	0.2	0.44	0.72	1.07	1.48	1.98


Outline

- 1 Motivation
- 2 Overview
- 3 ALSH-APPROX
- 4 MC-APPROX
- 5 Theoretical Analysis
- 6 Highlighted Experiments**
- 7 Final Remarks

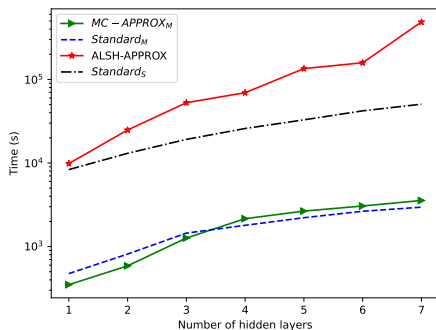


Time & Accuracy Comparison

Test accuracy




Training time




Time & Accuracy Comparison: MC-APPROX_s failed

MC-APPROX training time



Validation accuracies (7 hidden layers)

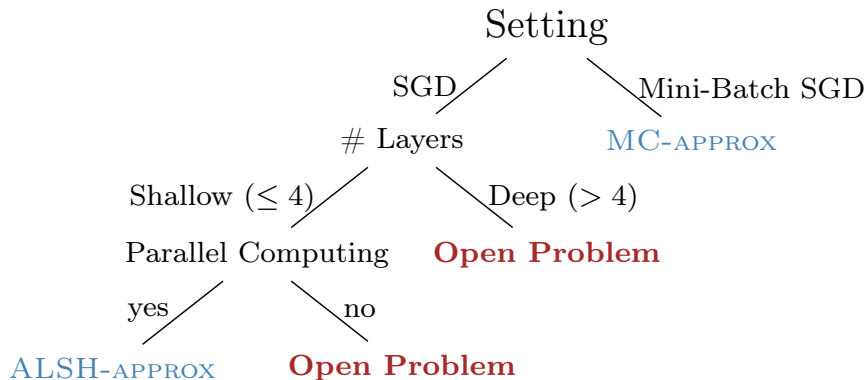


Outline

- 1 Motivation
- 2 Overview
- 3 ALSH-APPROX
- 4 MC-APPROX
- 5 Theoretical Analysis
- 6 Highlighted Experiments
- 7 Final Remarks**

Summary

Sampling-based Techniques for Training DNNs on CPU machines:



Next Step Advertisement!

Efficient Matrix Multiplication for Accelerating Inference in Binary and Ternary Neural Networks

- Complexity: $O(\frac{n^2}{\log n})$ time and memory
- Experiments: up to 6X less memory
- On CPU. C++: up to 30X faster; NumPy: 24X faster
- LLM Inference (Llama3.0, Falcon3, etc.), without system-level optimizations:
 - ▶ on CPU: up to 5.24X faster; on GPU: up to 2.5X faster



Thank you!

- InDeX Lab: cs.uic.edu/~indexlab/



Sana Ebrahimi



Rishi Advani

References



Menachem Adelman, Kfir Levy, Ido Hakimi, and Mark Silberstein.
Faster neural network training with approximate tensor operations.
In *NeurIPS*, 2021.



Petros Drineas, Ravi Kannan, and Michael W. Mahoney.
Fast Monte Carlo algorithms for matrices i: Approximating matrix
multiplication.
SIAM Journal on Computing, 36(1):132–157, January 2006.



Aristides Gionis, Piotr Indyk, and Rajeev Motwani.
Similarity search in high dimensions via hashing.
In *25th International Conference on Very Large Data Bases*, 1999.



Christian P. Robert.
Monte Carlo Methods.
Wiley, August 2016.



Ryan Spring and Anshumali Shrivastava.
Scalable and sustainable deep learning via randomized hashing.
In *ACM SIGKDD*, pages 445–454, 2017.