

# **CS480:** **Database Systems**

Introduction and Overview

Abolfazl Asudeh, [asudeh@uic.edu](mailto:asudeh@uic.edu)

Aug. 2025





## Instructor & TA

### Instructor:

- Abolfazl Asudeh
- <https://www.cs.uic.edu/~asudeh/>
- Email address: asudeh@uic.edu
- Drop-In Office Hours (in-person): **Mondays, 12:30 PM to 2:30 PM**
- Office number: 5452 CDRLC

### TA:

- Mohsen Dehghankar
- <https://mohsendehghankar.github.io/>
- Email address: mdehgh2@uic.edu
- Office Hours: **Wednesdays, 3:30 PM to 5:30 PM**
- Office number: CDRLC 2404



## Course Information

*What (Is a Database)?*

*Who (Uses DB)?*

*Why (Should you care)?*





## What Do Databases Do?

Core Functionality: Online Transaction Processing (**OLTP**)

- “**Define**” Data and their relationship
- “**Manipulate**” Data
  - **Insert**
  - **Update**
  - **Delete**



# What Do Databases Do?<sup>1</sup>

- Provide **persistent storage**
- Efficient **declarative** access to data (aka querying)
- **Atomicity** of updates: Protection from hardware/software **failures**
- **Concurrent** access by **multiple users**
- **Secure** and **Safe** access to data

---

<sup>1</sup> Some of the slides, including this one, are taken from Prof. Glavic's slides.



## What Do Databases Do?

### ACID Property

- **Atomicity:** Ensures that a transaction is treated as a single, indivisible unit of work. **All parts** of the transaction must complete successfully, or the entire transaction is rolled back, leaving the database in its original state.
- **Consistency:** Guarantees that a transaction brings the database from one valid state to another.
- **Isolation:** Ensures that the ongoing operations of one transaction do not interfere with other concurrently running transactions.
- **Durability:** Guarantees that once a transaction has been committed, its changes are permanent and will survive any subsequent system failures, such as power outages



## Brainstorming/curiosity question

**How would you do it?**



## Who Uses Databases?

- **Most big software systems involve DBs!**
  - Web-based systems
  - Business Intelligence; e.g., IBM Cognos
  - ...
- **Desktop / Mobile software (You!)**
  - Your music player; e.g., Amarok
  - Your email client
  - Half of the apps on your phone use SQLite
  - ...
- **Every big organization**
  - Banks, Google, ...
  - Government
  - ...



# Who Produces Database Systems?

- **Traditional relational database systems is big business**
  - IBM → DB2
  - Oracle → Oracle
  - Microsoft → SQLServer
  - Open Source → MySQL, Postgres, SQLite, DuckDB, . . .
- **Distributed systems with DB characteristics**
  - Cloud storage and Key-value stores → Amazon S3, Google Big Table, . . .
  - Big Data Analytics → Hadoop, Google Map&Reduce, . . .
  - SQL on Distributed Platforms → Spark, Tenzing, . . .



## Why Should You Care?

**49 Years of Queries** + 2 years

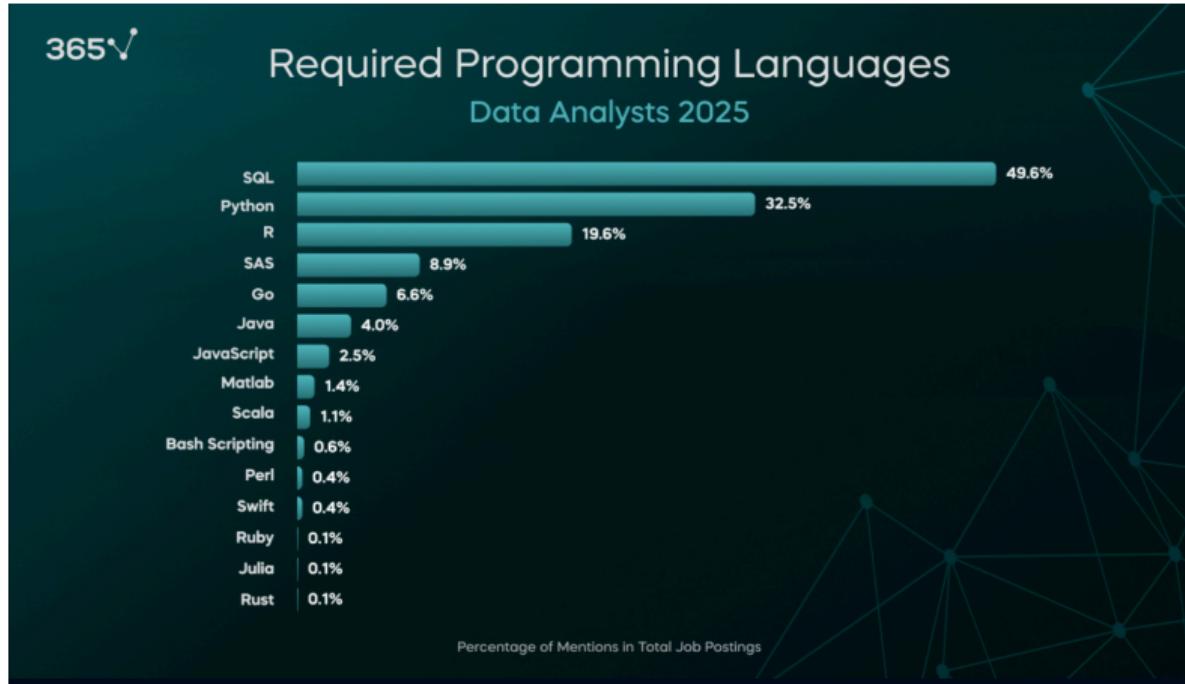
**Don Chamberlin**, SIGMOD Keynote 2023

**Abstract:** “The relational data model, proposed by **Ted Codd** in the 1970s, has been **the dominant paradigm** for storing and accessing business data for *several decades*.”





# 51 years of SQL, still at the top! (365datascience.com)





## Course Information

*Course Description*

*Outline of topics*

*Textbook*





## Course Description

This course has been updated to include the *recent trends and technological advancements* in AI and Computer Science. It is organized into two major parts.

- Part I: the core principles underpinning the discipline of database design and utilization.
- Part II: modern extensions of the course in the AI era.



## Outline of topics

Part I: **Relational Database Core** (2/3rd of the semester).

- Introduction to Databases & DBMS
- Data Modeling with ER Diagrams
- Relational Model and Constraints
- SQL: Basic Queries & Schema Definition
- SQL: Advanced Queries
- Functional Dependencies & Normalization



## Outline of topics

Part II: **Modern Extensions** (1/3rd of the semester).

- Vector Databases & Embeddings
- RAG: Retrieval Augmented Generation
- Fundamentals of NL-to-SQL with LLMs



## Excluded (yet important) topics – beyond the scope of this course

- **Relational Algebra**
- **System & Performance**
  - Indexing Structures & Query Processing
  - Transactions & Concurrency Control
  - Physical Database Design & Storage
- ...



## Course Materials: Part I: Relational Database Core

1. (main textbook: <https://db-book.com/>)  
Silberschatz, Korth, and Sudarshan, **Database System Concepts**, McGraw Hill, 2010, ISBN-13: 978-1260084504
2. Elmasri and Navathe, **Fundamentals of Database Systems**, 7th edition, Addison-Wesley, 2006, ISBN: 9780133971224
3. Ramakrishnan and Gehrke, Database Management Systems, 3rd Edition, McGraw-Hill, 2002, ISBN-13: 978-0072465631
4. Garcia-Molina, Ullman, and Widom, **Database Systems: The Complete Book**, 2nd Edition, Prentice Hall, 2008, ISBN-13: 978-0131873254



## Course Materials: Part II: Modern Extensions

### 1. Vector Databases:

- Online materials for Vector Database
- pgvector documentations and tutorials:
  - GitHub (official documentation): <https://github.com/pgvector/pgvector>
  - “PostgreSQL as a Vector Database: A Pgvector Tutorial”
- HNSW: Efficient and robust approximate nearest neighbor search using hierarchical navigable small world graphs.
- HENN: A Hierarchical Epsilon Net Navigation Graph for Approximate Nearest Neighbor Search. (The TA's paper)

### 2. RAG: TBA

### 3. NL-to-SQL: (a recent survey) Liu, Xinyu, et al. "A Survey of Text-to-SQL in the Era of LLMs: Where are we, and where are we going?." IEEE Transactions on Knowledge and Data Engineering (2025)



## Performance Evaluation

*Grading Policy*

*Points Breakdown*

*Project*





## Grading Policy

<b>1. Assignments</b>	$15\% = 5\% + 5\% + 5$	<ul style="list-style-type: none"><li>• A: &gt; 85</li></ul>
<b>2. Midterm</b>	30%	<ul style="list-style-type: none"><li>• B: &gt; 70</li></ul>
<b>3. Project</b>	20%	<ul style="list-style-type: none"><li>• C: &gt; 55</li></ul>
<b>4. Final exam</b>	35%	<ul style="list-style-type: none"><li>• D: &gt; 40</li><li>• F: &lt; 40</li></ul>



## Academic Honesty & Late Submission Policy

Cheating will NOT be tolerated!

- **All work has to be original.**
- **Cheating = 0 points for assignment/exam.** Possibly F in the course and **further administrative sanctions.**
- Every dishonesty **will be reported** to office of academic dishonesty
- Unless explicitly specified, you are **NOT ALLOWED** to use Generative AI and Large Language Models (such as ChatGPT).

Late policy for assignments and the project:

- **-20% per day**
- No exceptions!



# Project

- Create **groups of at most three**
  - Your responsibility!
  - Inform me + TA
- **Git repositories**
  - Create an account on GitHub, using your UIC email
  - Use it to exchange code with your fellow group members
  - The project has to be submitted via the group repository
- More details TBA!



## Project

- Every student has to contribute to every phase of the project!
- Don't let others freeload on your hard work!
- **Inform me or a TA ASAP**



## Software

*Git*

*Docker*

*PostgreSQL*





## Self Study

I expect you to learn by yourself how to effectively use the following technologies:

- **Git** - a version control system
  - Projects are submitted through git
  - We provide some useful examples/scripts through git
- **Docker** - a virtualization platform (think VMs, but more lightweight)
  - The easiest way to get Postgres running is by using the Docker image provided by Dr. Glavic's Github Repo (link below)
- **PostgreSQL**
  - I expect you to learn how to start/stop/configure a Postgres server and how to connect to a running Postgres server

### Further Instruction

Follow <https://github.com/lordpretzel/cs480>



# Postgres Overview

- Client/Server Architecture
- Postgres Cluster
  - A directory on the machine running the server that stores data and configuration files
- Postgres Server
  - A Postgres server handles the data of a single cluster
  - Clients connect to the server via network (TCP/IP)
    - Send commands and receive results



# Postgres Overview

- Clients
  - GUI clients: e.g., PGAdmin (<https://www.pgadmin.org/>)
  - CLI clients: e.g., the built-in psql tool
  - Programming Language Libraries
    - Java: JDBC (<https://jdbc.postgresql.org/>)
    - Python: asyncpg (<https://github.com/MagicStack/asyncpg>), psycopg (<http://initd.org/psycopg/>)



## Get Your Hands Dirty

Get a working version of the PostgreSQL server. Feel free to use LLMs (like ChatGPT) as a help for installing Postgres and getting it to work. Your options:

- **Install locally**
  - Installer packages for Windows exists
  - Linux distributions have a Postgres package
  - Installation from source is not that hard
- **Get our docker image** (docker pull Dr. Glavic's Github Repo)
  - It's an extension of the official Postgres image which loads our running example university database



## Get Your Hands Dirty

- Create a database cluster (the directory PostgreSQL uses to store data)
- Check that you can start/stop the server
- Check that you can connect to the running server using psql or any other client



## An AI-assisted tool for ERD design

Dr. Rooshenas has created a tool for designing ERD, which provides an AI assistant. You can use this for designing your ERD and ensuring its correctness.

Follow <http://aiassist.cs.uic.edu/>.

How to use: [link]

You will need a user name to use this – will be done soon.



Thank you!